

SUMMARY

AI Engineer with an M.S. in Data Science and 2+ years of professional experience building AI/ML applications, data pipelines, and analytical systems. Build LLM-backed systems across ingestion, embeddings, retrieval, context construction, grounded generation, evaluation, APIs, testing, and cloud delivery. Hands-on with RAG, hybrid retrieval, prompt/context engineering, custom agent orchestration, tool routing, structured outputs, validation, fallbacks, and human-review checkpoints.

TECHNICAL SKILLS

LLMs & Models	GPT, Claude, Gemini, Llama, Mistral, DeepSeek, Qwen
Agentic AI & Platforms	Groq, OpenRouter, Ollama, Hugging Face, LM Studio, prompt engineering, context engineering, agentic workflows, multi-agent workflows, custom agent orchestration, tool calling/routing, structured outputs, model routing, confidence thresholds, validation, fallbacks, human-in-the-loop workflows, multimodal AI
Retrieval & RAG	RAG, embeddings, semantic retrieval, hybrid retrieval, query rewriting, reranking, thread-aware context construction, grounded generation, Sentence Transformers, ChromaDB, FAISS, SQLite FTS
Backend & Delivery	Python, FastAPI, Flask, REST APIs, Docker, AWS/cloud, API integration, async workflows, caching, inference serving, Git, Pytest, Postman, integration testing, troubleshooting
Machine Learning & Data	Pandas, NumPy, SQL, SQLite, MySQL, scikit-learn, TensorFlow/Keras, transfer learning, ensemble/time-series models, SHAP, DoWhy, EconML, ETL, data validation, Power BI, Tableau

EXPERIENCE

AI Engineer Warren and Carter	Feb 2026 – Present
- Build end-to-end LLM/RAG workflows covering document ingestion, preprocessing, embeddings, semantic retrieval, context construction, grounded inference, and API-served delivery. - Design custom multi-step agent workflows with tool routing, structured outputs, validation rules, fallback paths, and human-review checkpoints; define evaluation criteria and failure modes for reliability-sensitive AI behavior. - Develop Python AI services with FastAPI/Flask and REST APIs, package components with Docker, and use AWS/cloud for selected deployments and backend integrations. - Test AI/API behavior with Pytest and Postman, troubleshoot integration failures, define observable outputs, and evaluate multi-modal text-and-vision approaches against product requirements.	
Data Analyst Tubman Technologies Inc.	Feb 2025 – Feb 2026
- Built Python, SQL, and Pandas ETL pipelines across healthcare and financial datasets, performing transformation, reconciliation, and validation before downstream analytics and AI use. - Prepared structured and document data for embedding and retrieval workflows through cleaning, normalization, and quality checks, creating dependable inputs for semantic-search and RAG systems. - Delivered Power BI/Tableau analyses of data-quality and utilization trends, translating technical findings into actionable insights for technical and non-technical stakeholders.	
Data Analyst SYMNN	Jul 2022 – Jul 2023
Analyzed telecom/network-utilization data and built predictive, ensemble, and time-series models in Python/scikit-learn to forecast utilization and support capacity-planning decisions.	

SELECTED AI / ML PROJECTS

Mailayer – Gmail Intelligence & Grounded AI	Python / FastAPI / Gmail API / SQLite/FTS / LM Studio / React/Vite
- Built a Gmail intelligence system with read-only OAuth, full/incremental mailbox synchronization, lexical and embedding-based retrieval, hybrid search, query rewriting, reranking, and thread-aware context construction. - Built custom retrieval and LLM orchestration for mailbox-aware question answering, grounding responses in retrieved messages and returning source emails for user verification. - Implemented priority scoring, classifier confidence thresholds, model routing, fallback behavior, asynchronous processing, caching, and user-feedback loops across retrieval and inference workflows. - Built custom retrieval and grounding evaluation cases to compare ranking behavior and identify weak evidence selection during system iteration.	
DRAX TBS – Document RAG Tutoring System	FastAPI / PyMuPDF / Sentence Transformers / ChromaDB
- Built an end-to-end document RAG pipeline spanning PDF ingestion, Sentence Transformer embeddings, ChromaDB vector retrieval, context assembly, and configurable local-LLM inference through LM Studio. - Implemented document, upload, chat, and session APIs with SQLite persistence and backend test support, separating ingestion, retrieval, model-serving, persistence, and API concerns.	
DaChat – CSV Chat & ML Predictions	Python / Streamlit / Pandas / scikit-learn / LM Studio
Built a Streamlit app for natural-language CSV querying and AI-powered property-price predictions, integrating LM Studio with trained scikit-learn model artifacts.	
Master’s Project – Facial Emotion Detection	VGG16 / ResNet50 / DenseNet121 / Transfer Learning
Developed and benchmarked four transfer-learning CNN architectures for multi-class facial-emotion recognition using preprocessing, augmentation, and class-imbalance handling; VGG16 achieved 97.83% test accuracy against ResNet50 and DenseNet121-based alternatives.	

EDUCATION

M.S. Data Science – New Jersey Institute of Technology (NJIT) Rel. Coursework: ML, AI, Deep Learning, RL, Applied Statistics, Big Data	Sep 2023 – Dec 2024
B.S Computer Science & Engineering – GITAM, Hyderabad Rel. Coursework: Big Data, AI, Data Warehousing & Mining, Robotics & Automation, Cloud Computing, DBMS	Jun 2019 – May 2023